

Orchestration isn't new. It's the problem OPERATE™ was built to solve.

A closer look at what the multi-model AI shift means for governance and trust.

THE SHIFT

The individual AI model is no longer the product.

90%

of enterprise tokens predicted to shift to open-weight models within 12 months.

— Matan Grinberg, CEO, Factory

70%

of engineering teams will route tasks through AI gateways by 2028, up from 25% in 2025.

— Gartner

THE REAL PROBLEM

This isn't a technology story.

When orchestration routes tasks across different models, users get inconsistent tone, reasoning, and behavior — from what looks like one tool.

Analysts call it the nondeterminism problem.

**It's not a training gap.
It's a legibility gap.**

DoorDash built DashBench.

An internal evaluation framework testing AI code review against their own historical pull requests, instead of public coding benchmarks.

Public benchmarks routinely fail to predict real-world performance.

THE NUMBERS

Three tiers. One benchmark.

No-scout baseline	30.7%
DoorDash's existing production system	53.6%
DashBench-optimized pairing	65.2% / 75.3%
<i>recall / F1, at \$3.81 per review</i>	

THE TAKEAWAY

The finding isn't that orchestration beats a single model.

It's that internal evaluation discipline found a better system than the orchestration DoorDash had already deployed.

Public benchmarks would not have surfaced that.

THE GOVERNANCE GAP

78%

of organizations use AI in at least one function.

18%

have enterprise-wide governance councils for a multi-vendor model ecosystem.

— McKinsey

The EU AI Act's high-risk system obligations become fully enforceable in August 2026.

Three stages. Already built for this.

01

Responsible AI Governance

Risk-tiered architecture, named human owners, identity access management for agents.

02

Trust Engineering

Provider-side behavioral drift, black-box liability, explainability.

03

Embed, Scale & Evolve

Token-level cost tracking, LLM-as-a-judge evaluation frameworks.

This is the discipline the orchestration shift now confirms enterprises need.

The ADOPT loop runs alongside ADKAR.

The Institute of Project Management proposed ADOPT, Awareness, Define, Oversee, Prompt, Trust, as a separate operating model for supervised AI.

OPERATE takes a different approach.

Governance and trust discipline built as native stages within a framework designed to supplement ADKAR, not run parallel to it.

THE QUESTION

Does anyone in your organization hold named, accountable ownership for what your AI system does when it routes a task to a model nobody explicitly chose?

Full playbook: link in comments.

Selvakumar Jayakrishnan

Senior Change & Transformation Leader | 19 years enterprise, 11 years core Change Management

selvakumarjayakrishnan.com | linkedin.com/in/passionateselva